




PLAYBOOK · EDITION 01

The Ocean Freight Automation Playbook

14 workflows for the forwarders who ship before their competitors reply.

A NOTE FROM THE TEAM

We built Susea because we watched too many forwarders lose deals they should have won.

Not because their rates were wrong, but because they answered four hours late.

We spent the first six months of our beta not writing code. We sat next to operators in Mumbai, Rotterdam, and Dubai. We watched them open Outlook at 08:42, count 47 unread RFQs, and try to figure out which three mattered before their coffee went cold.

The workflows in this playbook are the ones those operators asked us for. Some of them we built into Susea. Some of them we didn't — but they're worth knowing about anyway, because a pattern is a pattern whether the system running it is Susea, a scrappy Zapier setup, or a genuinely well-trained ops lead.

We labelled the shortcuts honestly. Where Susea does something out of the box, we say so. Where it doesn't, we tell you what our beta customers hacked together instead.

Steal freely.

— The Susea team

THE 14 AT A GLANCE

Fourteen workflows, three groupings.

REQUESTS COMING IN

- 01 RFQ intake — Turning five channels into one queue. p.4
- 02 Tariff intelligence — Catching rate changes before they hurt. p.6
- 03 Customer profile enrichment — Quote from context, not from scratch. p.8

QUOTES GOING OUT

- 04 The all-in rate draft — Under 90 seconds, deterministic math, transparent line items. p.10
- 05 Margin-aware approval routing — Who says yes at what threshold. p.12
- 06 Multi-currency & FX drift handling — Quote in the customer's currency without giving away margin overnight. p.14

AGENTS & AWARDS

- 07 The silent-quote recovery cadence — 32% of quiet quotes are recoverable. Most forwarders never send a second touch. p.16
- 08 Quote expiry re-price — Refresh the rate before the customer has to ask. p.18
- 09 Won / lost signal capture — If you don't know why you lost, you'll lose it again next quarter. p.20
- 10 RFQ broadcast & response normalisation — Send a lane to six carriers. Get back one comparable table. p.22
- 11 Bid chasing — Nudge the non-bidders. With time-left and ranking context, not spam. p.24
- 12 Bid QA — Catch bad bids before they reach your award screen. p.26
- 13 Award justification & audit trail — Every award needs a reason. Every reason needs a paper trail. p.28
- 14 Contract utilisation alerting — The negotiated saving isn't real until you actually ship the volume. p.30

01

REQUESTS COMING IN


RFQ intake

Turning five channels into one queue.

Requests arrive on email, WhatsApp, portal, phone, and walk-in. Whoever sees an RFQ first “owns” it — until they don’t. Duplicates get quoted twice. Others sit in an inbox nobody’s reading. The customer who asked three forwarders picks whoever answered first — and that’s rarely the forwarder whose ops manager was on a call.

The workflow that fixes this doesn’t try to force customers into one channel (they won’t). It forwards every channel into a single intake surface, parses the incoming request the same way regardless of source, dedupes against open records, and lands a structured RFQ in a triage queue where it’s impossible to miss.

THE WORKFLOW



TRIGGER LOGIC

- Any inbound where the message contains freight-request signal: at least one of POL, POD, container type, cargo description, sailing date.
- Trigger fires within 60 seconds of message arrival — the whole point is time-to-triage.

EDGE CASES TO BUILD FOR

- **Multi-lane RFQs** (one email, six lanes) — parser must create N structured records, not one.
- **Unknown sender** — route to an “unmatched sender” queue with a one-click “add to CRM” action.
- **Incomplete info** — mark missing fields explicitly rather than guessing.
- **Mixed-language requests** (English + Hindi / Arabic / Spanish) — parser must handle bilingual messages without failing silently.
- **Follow-up on an existing RFQ** — should update the record, not create a duplicate.

DON'T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the “obvious” trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- 100% capture rate — no RFQ arrives without an operator seeing it within 15 minutes.
- Zero duplicate quotes to the same customer within 24 hours.
- Time-to-triage under 5 minutes.

THE SUSEA SHORTCUT

Susea's **a-01 · AI quotation drafting** does the parsing, deduping, and structured-record creation across email and WhatsApp out of the box. The unmatched-sender queue is built in. If you'd rather build it yourself: n8n or Make.com + an LLM extraction step + a CRM webhook will get you 70% of the way there, but plan for the mixed-language edge case early.

02

REQUESTS COMING IN


Tariff intelligence

Catching rate changes before they hurt.

A carrier posts a GRI on Tuesday. Nobody on your desk sees it until Friday morning. Between Tuesday and Friday, three quotes go out on the old rate. Two of them get accepted. You just ate 4.2% of margin on those shipments and nobody knows.

The fix isn't a smarter operator. It's a workflow that treats every carrier communication — email, PDF attachment, WhatsApp forward — as a potential rate change, extracts the delta, and flags every open quote it touches before someone hits send on a stale price.

THE WORKFLOW



TRIGGER LOGIC

- Any inbound with rate-change signal (“GRI,” “PSS,” “rate update,” “effective from”), any channel.
- Also: scheduled sweep of carrier portals if credentials are available.

EDGE CASES TO BUILD FOR

- **Mid-month partial GRIs** (some lanes, not others) — extraction must respect scope, not apply blanket updates.
- **Region-specific changes** — an Asia-EU rate change shouldn’t touch Asia-US quotes.
- **Private / agent-negotiated rates** that override public tariff must not be overwritten by public rate updates.
- **GRIs announced then withdrawn** — the workflow needs a revert path.

DON'T DO THIS

[Team to fill in: 2–3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the “obvious” trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- Under 2 hours from carrier posting a change to your pricing lead seeing it.
- Zero quotes sent on rates superseded more than 4 hours before send.
- Tariff database version-controlled: you can always answer “what was our rate on this lane on X date?”

THE SUSEA SHORTCUT

*Two Susea automations combine here: **a-03 · Rate & GRI alerts** watches carrier communications and flags open quotes affected; **a-06 · Tariff intake** handles the actual database update from PDF / XLS / email / WhatsApp forwards. Together they make Tuesday’s GRI a Tuesday event, not a Friday surprise.*

03

REQUESTS COMING IN


Customer profile enrichment

Quote from context, not from scratch.

Every quote starts blank. If the operator remembers that Acme prefers CFR, has a 3.5% margin band, ships mostly out of Nhava Sheva, and needs their quotes in EUR — the quote is faster and better. If they don't — and rotating desks mean they often don't — the quote is generic, sometimes wrong, and always slower.

The fix is a customer profile that carries context across quotes: default incoterm, currency, margin band, preferred sailings, historical lanes, do-not-quote flags, and past acceptance patterns. It's not a feature — it's a data model. And it makes every downstream workflow smarter.

THE WORKFLOW



TRIGGER LOGIC

- New RFQ where sender email domain or company name matches an existing customer.
- Also: manual “assign to customer” from the operator when the match isn’t automatic.

EDGE CASES TO BUILD FOR

- **Multiple contacts per customer** with different preferences (procurement wants EUR, ops wants USD).
- **Subsidiaries and holding companies** — same commercial terms may or may not apply.
- **Walk-ins who become repeat customers** — profile must be creatable retroactively without losing history.
- **Do-not-automate flag** for accounts under commercial dispute or renegotiation.

DON'T DO THIS

[Team to fill in: 2–3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the “obvious” trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- Quote-fill time on repeat customers cut by 40%.
- Operator adjustments visible in the draft — nothing hidden.
- Even a first-day operator on the desk quotes Acme with Acme’s preferences.

THE SUSEA SHORTCUT

Susea holds customer profile as a first-class object in the data model — every automation pulls from it. There’s no single automation labelled “customer profile” because it’s the substrate everything else runs on. If you’re building this yourself: even a Notion / Airtable customer table linked to your CRM will get you most of the way, as long as it’s the one your operators actually check.

04

QUOTES GOING OUT


The all-in rate draft

Under 90 seconds, deterministic math, transparent line items.

An RFQ comes in. Your operator opens the tariff (which version?), calculates ocean freight, adds BAF, THC, ISPS, LSS, documentation, applies the margin — then formats the email. Eighteen minutes on a good day. On a Monday morning with 47 in the queue, it never happens for most of them.

The workflow that fixes this is not “AI-generated quotes.” It’s deterministic quote assembly, with AI only reading the input and formatting the output. The math stays yours — tariff, surcharge stack, margin band. The speed comes from removing every step that used to be manual retyping.

THE WORKFLOW



TRIGGER LOGIC

- Structured RFQ has minimum viable info: lane, container type, sailing window.
- If any is missing, workflow doesn't draft — it triggers Workflow 01's "missing info" clarification instead.

EDGE CASES TO BUILD FOR

- **Split shipments** (partial loads across sailings) — line items multiply.
- **Multi-container mixed types** (2×40'HC + 1×20') — surcharge stack applied per container type.
- **Port-pair with no direct service** — draft must offer transshipment routing with clear disclosure.
- **Sailing windows outside quoted validity** — flag before send, not after.

DON'T DO THIS

[Team to fill in: 2–3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the "obvious" trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- Median draft time under 90 seconds.
- Zero cases where AI touched the price arithmetic.
- Every line item traceable to a specific tariff row and customer-profile setting.

THE SUSEA SHORTCUT

*This is Susea's flagship — **a-01 · AI quotation drafting** does the end-to-end draft under 90 seconds by default. The critical design choice: AI drafts the wrapper, deterministic logic runs the math. If you're building this yourself, keep that separation religiously — the moment you let the language model do the surcharge arithmetic, you've lost the audit trail your finance team needs.*

05

QUOTES GOING OUT


Margin-aware approval routing

Who says yes at what threshold.

If every quote goes to the same approver, that person becomes the bottleneck at scale. If nothing gets approved and operators self-approve to keep customers happy, margins drift down quietly. Neither is fine.

The workflow that fixes this treats approval as a rule, not a habit. Quotes above your margin threshold auto-approve. Quotes below threshold route to a named approver with an SLA clock. Strategic accounts and first-time customers can override the rule. The whole thing is auditable.

THE WORKFLOW



TRIGGER LOGIC

- Every quote draft, before send. Non-negotiable.
- Threshold configurable per lane, customer band, and product (spot vs. contract).

EDGE CASES TO BUILD FOR

- **Strategic account override** — allow named accounts to bypass threshold with reason logged.
- **First-time customer** — apply stricter threshold until relationship data exists.
- **Contract renewal quotes** — different threshold than net-new.
- **SLA breach** — if approver doesn't respond within X hours, escalate up, not down.

DON'T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the "obvious" trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- Approver bottleneck reduced by more than half.
- Zero cases of below-threshold quotes self-approved without a paper trail.
- Approvers report they see fewer, higher-signal decisions.

THE SUSEA SHORTCUT

*Susea's **a-05 · Auto-approval routing** handles the threshold check, routing, and audit log natively. Thresholds are per-lane and per-customer-band. If you're building this yourself: any workflow tool + your CRM's user roles can approximate it, but the audit log is where most homegrown versions fall down — that's the one thing you can't skip.*

06

QUOTES GOING OUT

Multi-currency & FX drift handling

Quote in the customer's currency without giving away margin overnight.

Your customer wants a EUR quote. Your tariff is USD. Between the moment you draft the quote and the moment the customer accepts it, FX moves 1.5%. Whose margin absorbed the shift? If you don't have a workflow for this, the answer is: yours, and you don't know it.

The workflow that fixes this locks the FX rate at quote draft time, tracks drift while the quote sits open, and re-quotes with fresh FX if the drift crosses your tolerance — before the customer sees a surprise.

THE WORKFLOW

TRIGGER**Quote in non-tariff currency**

Customer wants EUR, tariff USD

DETERMINISTIC**Lock FX rate**

At draft timestamp

SCHEDULED**Check drift**

Every 12h while quote is open

DECISION**Drift > tolerance?**

Default 0.75%

HUMAN**Operator sees alert**

Re-quote drafted

OUTPUT**Re-quote sent**

Framed as "updated FX"

TRIGGER LOGIC

- Every quote in a currency different from your primary tariff currency.
- Drift check runs on a schedule (12h default; hourly for high-volatility currencies).

EDGE CASES TO BUILD FOR

- **Customers who want both currencies visible** (USD reference + local currency actual).
- **INR-USD and BRL-USD volatility spikes** — tighter drift tolerance for these pairs.
- **Customer accepts at locked rate** after drift alert but before re-quote is sent — honour the locked rate.
- **FX rate source unavailable** temporarily — hold, don't guess.

DON'T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the "obvious" trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- FX-driven margin surprise under 0.5% across the quarter.
- Customers experience currency handling as "proactive," not "flaky."
- Finance can point to the exact FX rate used on every historical quote.

THE SUSEA SHORTCUT

Susea handles this inside the quote engine rather than as a labelled automation — the FX-lock and drift-tolerance settings live in your workspace config, and the drift alert routes through the same approval logic as any other margin-affecting change. If you're building this yourself, the drift check is easy — the hard part is the "customer accepted at the locked rate, honour it" logic. Design that first.

07

AGENTS & AWARDS

FLAGSHIP


The silent-quote recovery cadence

32% of quiet quotes are recoverable. Most forwarders never send a second touch.

You sent Q-2416 on Tuesday. The customer went quiet. On Friday afternoon, if someone remembers, a rep might chase it. Might not. Most quotes that go quiet on a Tuesday are never followed up on — the operator moved on to the next fire, and the quote quietly aged into “closed / no response.”

Our beta cohort’s data says roughly a third of quiet quotes come back if you send the right follow-up at the right time. The workflow that captures those wins is a cadence — timed, personalised, and stopping the moment the customer responds.

THE WORKFLOW



TRIGGER LOGIC

- Quote sent, no reply received on any channel (email, WhatsApp, phone log entry).
- Reply detection must span channels — a WhatsApp reply cancels an email follow-up.

EDGE CASES TO BUILD FOR

- **Customer replies on a different channel** than the quote was sent on (very common: quote emailed, reply on WhatsApp).
- **Weekend / holiday timing** — don't send a nudge Sunday evening.
- **Customer explicitly said "will get back next week"** — pause the cadence until then.
- **Automated out-of-office replies** — don't count as a reply.
- **Bounce / undeliverable** — flag the operator immediately.

DON'T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the "obvious" trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- 32% of quiet quotes recovered (beta cohort benchmark).
- Fewer than 5% of customers report follow-ups as "too many."
- Zero cases of a follow-up sent after the customer already accepted.

THE SUSEA SHORTCUT

Susea's a-02 · Silent-quote recovery runs this cadence end-to-end, with cross-channel reply detection built in. The 32% figure comes from Susea beta customers' actual recovery rates. If you're building this yourself: the cadence is the easy part. The cross-channel reply detection is what separates a workflow customers appreciate from one they resent.

08

AGENTS & AWARDS


Quote expiry re-price

Refresh the rate before the customer has to ask.

A quote valid till Friday. Customer replies on Monday. The rate has moved. Now you're either eating margin on your side, or sending a "sorry, revised rate" email that makes you look flaky.

The workflow that fixes this catches the quote before validity expires, checks the current rate, and either extends the quote (if rates are stable) or drafts a proactive re-quote (if they aren't) — framed as a service, not an apology.

THE WORKFLOW



TRIGGER LOGIC

- 24h before validity expires. Adjust to 48h for lanes with historically slow customer decisions.
- Only if quote hasn't been accepted / lost.

EDGE CASES TO BUILD FOR

- **Customer verbally / informally accepted** but hasn't sent formal PO — don't re-price, chase for PO instead.
- **Contract-rate quotes** (as opposed to spot) — the whole workflow is skipped; contract terms hold.
- **Space tightening on the sailing** — even at stable rate, the extension may need a "subject to space" flag.
- **Multi-quote customer** (they got quotes on 5 lanes, replying only on 2) — re-price only the unaddressed ones.

DON'T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the "obvious" trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- Zero "sorry, revised rate" emails sent reactively.
- Customers experience the re-quote as service, not backpedalling.
- Every re-quote clearly cites what changed and why.

THE SUSEA SHORTCUT

Susea's expiry re-price runs proactively for every open quote — the automation is listed as one of the eight on susea.ai/automations. The workflow shares logic with Silent-quote recovery (Workflow 07), which is why they route into each other. If you're building this yourself: the "extend vs. re-price" decision is where most homegrown versions get it wrong — err on the side of transparency with the customer, always.

09

AGENTS & AWARDS


Won / lost signal capture

If you don't know why you lost, you'll lose it again next quarter.

You quote 100 things this week. 30 win. 40 lose. 30 go quiet and eventually age out. Ask your commercial lead why the 40 lost, and you'll get anecdotes, not answers. Next quarter you'll make the same pattern of mistakes — because nobody wrote them down at the moment they happened.

The workflow that fixes this captures a close signal at the exact moment a quote resolves, structured enough to analyse later. Not every quote — that's too much. But enough that patterns become visible: which lanes you keep losing on, which customer segments you're mispriced for, which competitor keeps beating you.

THE WORKFLOW



TRIGGER LOGIC

- **Won:** customer accepts (email, WhatsApp, formal PO) — any explicit acceptance.
- **Lost:** customer says no explicitly, or awarded to a competitor visibly.
- **Aged out:** 14 days silent past validity expiry with no follow-up response.

EDGE CASES TO BUILD FOR

- **Won on unrelated factors** (relationship, timing, favour) — capture that context, don't overweight the pricing signal.
- **Lost to a customer who wasn't real** (broker fishing for benchmarks) — mark and exclude from pricing analytics.
- **Delayed decisions** ("we'll do this next quarter") — separate from lost, track as pipeline.
- **Split awards** — customer took 2 of your 3 lanes, gave 1 to competitor. Capture the split.

DON'T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the "obvious" trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- Signal captured on 70%+ of resolved quotes.
- Loss reason patterns visible within a quarter.
- Pricing team uses the data monthly, not "someday."

THE SUSEA SHORTCUT

Susea holds won / lost as first-class fields on the quote object — no separate automation, it's part of the data model. The value compounds: the more quotes resolve with signal, the smarter every downstream workflow gets. If you're building this yourself: fight the temptation to make loss-reason mandatory. Optional-but-prompted gets you higher-quality data than forced clicks.

10

AGENTS & AWARDS


RFQ broadcast & response normalisation

Send a lane to six carriers. Get back one comparable table.

You send a lane to six carriers. Come back: two PDFs, one Excel, one WhatsApp voice note, two emails with prose. Now your ops lead spends the next hour reformatting six responses into one comparable table before anyone can decide.

The workflow that fixes this treats every response — regardless of format — as data that needs to hit the same schema. Extract, normalise, table it. The comparison is where value is; the reformatting is where an hour disappears.

THE WORKFLOW



TRIGGER LOGIC

- RFQ record marked “broadcast.” Recipient list can be internal (agents) or external (carriers directly).
- Reminder to non-responders at 50% of RFQ window elapsed.

EDGE CASES TO BUILD FOR

- **Carriers who won’t quote without volume commitment** — separate “no bid, needs volume signal” from “no bid, no reason.”
- **Agent bidding in unusual currency** — normalisation converts, but original currency preserved on the record.
- **Split-quote responses** (they quote 1 of 3 lanes) — must capture partial responses cleanly.
- **Off-schedule reply after award already made** — log, don’t insert into the comparison retroactively.

DON’T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the “obvious” trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- Under 2 hours from broadcast to comparable table.
- 100% of responses captured — nothing lost in an operator’s inbox.
- Comparison shows rate, transit, and reliability side-by-side, not just rate.

THE SUSEA SHORTCUT

*Susea’s **a-04 · RFQ broadcast** handles the send, the receive, and the normalisation. Responses arrive on any channel and land in one table. If you’re building this yourself: the normalisation is the hard part. Reception is easy (a shared inbox); making six different response formats comparable is where the workflow earns its keep.*

11

AGENTS & AWARDS


Bid chasing

Nudge the non-bidders. With time-left and ranking context, not spam.

You launched an RFQ. Deadline is Thursday. It's Wednesday morning. Half your agents haven't bid yet. Do you send a mass reminder email nobody reads, call each of them individually, or let the deadline pass and wonder why bid count was low?

The workflow that fixes this treats non-responders as data — who hasn't bid, how much time is left, where they currently rank against the bids already in. AI drafts a personalised nudge for each agent with that context. You approve, it sends. Nobody feels chased; everyone knows the clock is ticking.

THE WORKFLOW



TRIGGER LOGIC

- 24h before RFQ deadline — every agent invited who hasn't bid yet.
- Optional earlier nudge at 50% of RFQ window elapsed, softer tone.

EDGE CASES TO BUILD FOR

- **Agent explicitly declined to bid** — don't chase them again.
- **Agent bid then withdrew** — different message than never-responded.
- **Partial bid** (some lanes not others) — nudge only on missing lanes.
- **Timezone mismatch** — don't send at 3am agent-local. Use business-hours window.
- **Agent OOO / auto-reply detected** — pause cadence, don't queue up nudges.

DON'T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the "obvious" trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- 90%+ bid response rate by RFQ close.
- Zero agents chased outside business hours.
- Agents don't complain about spam — nudge quality is the whole game.

THE SUSEA SHORTCUT

*Susea's **Bid chasing** automation runs this cadence natively — non-responder detection, per-agent context (time left + current ranking), personalised drafting, timezone and OOO guardrails. You approve; it sends. If you're building this yourself: the tempo and personalisation of nudges is where most homegrown versions fail. One thoughtful nudge with real context outperforms three generic reminders every single time.*

12

AGENTS & AWARDS


Bid QA

Catch bad bids before they reach your award screen.

Six carriers bid on your RFQ. One quoted in EUR when you asked for USD. One forgot to include BAF. One quoted a 4-day transit that literally isn't possible on that lane. You catch two of these in review. The third makes it to your award screen and creates a mess — either you award to a broken bid, or you unwind mid-decision.

The workflow that fixes this checks every bid the moment it lands against a set of validation rules — currency, surcharge completeness, transit-time plausibility, missing fields. Bad bids get flagged before they ever contaminate your award decision. Clean bids move through untouched.

THE WORKFLOW



TRIGGER LOGIC

- Every bid received on any open RFQ, checked within 60 seconds of submission.
- Checks run on structure and completeness, not on price. The math stays yours.

EDGE CASES TO BUILD FOR

- **Multi-lane bid with issues on some lanes only** — validate per lane, not per submission.
- **Agent bids in obscure currency** — convert for comparison, don't reject the bid.
- **Transit time flagged as "impossible"** — may be a new service you don't know about; treat as flag-for-review, not auto-reject.
- **Agent explicitly noted an exception** ("no BAF this quarter, replaced with EBAF") — parser must honour these notes.
- **Bid submitted after deadline** — separate flag entirely, not a QA issue.

DON'T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the "obvious" trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- Zero bad bids reach your award screen.
- Clarification cycle under 24h from flag to fix.
- Agents don't feel over-policed — only real, actionable issues get flagged.

THE SUSEA SHORTCUT

*Susea's **Bid QA** automation runs every check the moment a bid lands, batches clarifications back to the agent, and never touches the pricing math — only structural checks (currency, completeness, transit-time plausibility). If you're building this yourself: start with the two highest-cost checks — currency and missing surcharges. Add transit-time validation later, once you have reliable lane baselines to check against.*

13

AGENTS & AWARDS


Award justification & audit trail

Every award needs a reason. Every reason needs a paper trail.

You awarded to Maersk. Two weeks later, your ops lead asks why not MSC (they were cheaper). You can't remember. Or worse: your compliance review flags "no rationale on file." Or your customer's procurement team asks for the audit trail.

The workflow that fixes this drafts the rationale at the moment of award — while the reasoning is still fresh — and files it with the contract. Not "someday" — automatically, before the award is even communicated to the winning carrier.

THE WORKFLOW



TRIGGER LOGIC

- Award decision — status change on the RFQ record to “awarded.”
- Rationale must be drafted before the award notification goes to the winning carrier.

EDGE CASES TO BUILD FOR

- **Strategic reasons not in the data** (“relationship,” “reciprocity”) — draft must accept and preserve these without overweighting the quantitative comparison.
- **Split awards across carriers** on one RFQ — one rationale per split.
- **Single-source situations** (only one carrier bid) — rationale still needs to exist, framed as “sole bidder.”
- **Award reversal** (winner backs out) — old rationale archived, new one drafted for replacement.

DON'T DO THIS

[Team to fill in: 2–3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the “obvious” trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- 100% of awards have rationale on file.
- Under 2 minutes from award click to rationale drafted.
- Compliance / procurement audits pass without a scramble.

THE SUSEA SHORTCUT

*Susea’s **Award justification** automation drafts the rationale automatically at the moment of award. It’s stored with the contract record, so it’s there when someone asks three months later. If you’re building this yourself: the audit trail is the whole point — a rationale nobody can find later isn’t a rationale.*

14

AGENTS & AWARDS


Contract utilisation alerting

The negotiated saving isn't real until you actually ship the volume.

You fought for 20% off on 500 TEU per quarter with MSC. It's mid-quarter. You're at 180 TEU. Nobody's watching. End of quarter: you missed the volume commitment, lost the discount retroactively, and lost your negotiating leverage for next round. The saving you fought for was never real.

The workflow that fixes this sweeps contract commitments weekly, projects your utilisation trajectory, and alerts you before the shortfall becomes irreversible — with time to route more volume, or renegotiate the commitment.

THE WORKFLOW



TRIGGER LOGIC

- Weekly sweep on all active contracts with defined volume commitment.
- Immediate alert if a single lane's projected utilisation drops below 60% of committed volume.

EDGE CASES TO BUILD FOR

- **Seasonal volume commitments** (Q1 lighter than Q3 by design) — projection must account for seasonality, not straight-line.
- **MQC vs. targets** — MQC breach has commercial consequences; target miss doesn't.
- **Multi-lane contracts with carry-over rules** — one lane over, another under, net compliance intact.
- **Contract renegotiation in-flight** — pause alerts, resume when new terms lock in.

DON'T DO THIS

[Team to fill in: 2-3 real observations from beta customers. Suggested slots: (1) what a forwarder tried that backfired, (2) the "obvious" trap most teams fall into, (3) the edge case that only shows up at scale.]

WHAT GOOD LOOKS LIKE

- Zero contracts end-of-term below commitment without prior action.
- Commercial team acts on alerts within 5 business days.
- Savings you negotiated actually show up in your P&L.

THE SUSEA SHORTCUT

*Susea's **Contract utilisation alerts** automation runs this sweep by default. Every active contract is tracked, projections update weekly, and the alert routes to the named commercial owner on the contract. If you're building this yourself: the weekly sweep is straightforward. The hard part is making the alert land with someone who has time to act on it — most homegrown versions send it to a shared inbox where it's ignored.*

THE META-LESSON

What all fourteen have in common.

If you read the workflows carefully, you noticed something: the AI does the reading, the drafting, the routing, the extracting. It never does the pricing math. It never sends without approval. It never decides an award.

That's not a limitation. That's the design.

The workflows that work in freight ops are the ones that leave commercial judgement and pricing arithmetic where they belong — with the humans who own the P&L. AI belongs in the plumbing: reading a WhatsApp message into a structured record, checking six carriers for a stale rate, drafting a follow-up in the operator's voice, watching a contract's utilisation curve.

Your team keeps the judgement, the automation removes the switching.

That's the operating system underneath modern freight forwarding. Whether you build it yourself, hack it together, or use Susea to skip the plumbing entirely — that principle is the one worth stealing.

IF YOU'D RATHER NOT BUILD THESE YOURSELF

Susea runs eight of these fourteen out of the box.

Susea is the AI operating system for modern freight forwarders. It runs eight of the workflows in this playbook as native automations, and touches most of the other six through the underlying data model. Every draft is human-approved before it sends. Pricing math stays deterministic and yours. Nothing gets ripped out of your existing stack.

Our beta cohort onboards in three to five business days. A founder walks your ops lead through connecting carriers, importing your tariff sheet, and turning on the first two automations. You send real quotes from day one.

We're small enough that a founder personally onboards every customer this quarter. Beta pricing is locked in for twenty-four months. Four onboarding slots are open this month.

PLAYBOOK → SUSEA MAPPING

01 RFQ intake	Native
02 Tariff intelligence	Native
03 Customer profile enrichment	Data model
04 All-in rate draft	Native
05 Margin-aware approval routing	Adjacent
06 Multi-currency & FX drift	Quote engine
07 Silent-quote recovery	Native
08 Quote expiry re-price	Native
09 Won / lost signal capture	Data model
10 RFQ broadcast & normalisation	Adjacent
11 Bid chasing	Native
12 Bid QA	Native
13 Award justification & audit	Native
14 Contract utilisation alerting	Native

[Book a live demo →](#)
[Request beta access](#)

The Ocean Freight Automation Playbook.

Steal freely. Share widely.

THE SITE

susea.ai
susea.ai/automations
susea.ai/freight-forwarding

GET IN TOUCH

info@alphabitssolutions.com
linkedin.com/company/susea-ai

EDITION

© 2026 AlphaBits Solutions
Steal-and-share license